



SAPIENZA
UNIVERSITÀ DI ROMA

LA SAPIENZA UNIVERSITÀ DI ROMA

DIPARTIMENTO DI INGEGNERIA INFORMATICA, AUTOMATICA E
GESTIONALE

TRAJECTORY ENCODING

Authors:

Leandro de Souza Rosa

November 6, 2023

1 Trajectory Learning

2014 - Sequence to Sequence Learning with Neural Networks

Name: seq2seq

[1]

2017 - Identifying Human Mobility via Trajectory Embeddings

[2] classifies users based on trajectory data. The problem is hard because there are many more trajectories than users.

Recurrent Neural Network (RNN) is used, and said to be good for classification when the number of labels is small. In particular uses a Long Short-Term Memory (LSTM) for processing sub-trajectories.

There is a location embedding, not sure how that is computed. But the trajectories are points on google maps, so there maybe be semantic information in there. The sequence of location embedding is passed onto the LSTM, not sure how they handle different trajectory lengths.

2018 - Self-Consistent Trajectory Autoencoder: Hierarchical Reinforcement Learning with Trajectory Embeddings

Code on: <https://github.com/wyndwarrior/Sectar>

[3] learns an embedding for trajectories with one encoder and two decoders: a state decoder to decode from the latent space back into trajectories, and a policy decoder, which generates the trajectory in the environment. As such the state decoder predict the trajectory of the policy. The encoder is used in a hierarchical Reinforcement Learning (RL) setup.

The state encoder and decoder are RNNs and the policy decoder is a feed-forward Neural Network (NN).

note: if the encoder is trained with trajectories from different tasks, the policy will be conditioned to each task, what is sort of parameterizing the policy to tasks.

In the paper, the policy has unknown dynamics, and hence the RL setup. Trajectories are continuous poses of joints over time. Tested in simulation.

2018 - Anomalous Trajectory Detection Using Recurrent Neural Network

Code on: <https://github.com/LeeSongt/ATD-RNN>

[4] proposes anomalous trajectory detection using RNN.

The trajectories are discretized, using a grid, and feed to a stacked RNN for learning the embedding, then a multi-layered perceptron and a soft-max layer detects if the trajectory is anomalous. The stacked RNN is made by feeding the hidden states of the previous to the next RNN.

The trajectories are padded in order to get trajectories of the same length.

LSTM and Gated Recurrent Unit (GRU) are two special types of RNNs are tested. GRUs seems to work better.

2018 - Deep Representation Learning for Trajectory Similarity Computation

Name: t2vec Code on: <https://github.com/boathit/t2vec>

[5] presents t2vec. A Deep Learning (DL) approach for trajectory similarity. States that using RNN is not a very good idea because you cannot reconstruct the trajectory and it fails to consider spatial proximity, which is inherited in trajectory data.

called in the paper as t2vec or seq2seq?

The approach is based on the encoder-decoder framework. Handling varying sampling rates is done by augmenting the training data creating sub-trajectories by sub-sampling and noise addition. They also propose a spatial-aware loss, and pre-train the ??cells?? and let them to be optimized during training.

Notes: The paper is very confusing. I do not really know that are the inputs and outputs or how the sequences are fed in the RNN inside the encoder.

2020 - Trembr: Exploring Road Networks for Trajectory Representation Learning

[6] presents `traj2vec`. The paper focus on trajectories on roads. It preprocesses the trajectories by projecting them in a road network and the trajectory is a sequence of road segments and travel time.

The `RNN` decoder is conditioned to the road network, and the training is made by optimising a loss for the trajectory and another for time.

Notes: Maybe the secret for velocities profiles is in the addition of time to the loss.

2019 - Computing Trajectory Similarity in Linear Time: A Generic Seed-Guided Neural Metric Learning Approach

Name: NeuTraj

Code on: <https://github.com/yaodi833/NeuTraj>

[7] proposes a method for accelerating trajectory similarity computation by sampling seeds of trajectories, computing their similarity, and approximating them with a neural metric.

States that `RNNs`, `LSTMs`, and `GRUs` can only model one sequence without considering the between-sequence correlation.

Does not consider time in the trajectory. Starts sampling from the trajectories and computes a distance matrix between the samples using a given trajectory distance metric which is then normalized.

The `RNN` is augmented with a memory, which is created by dividing the space into a grid, and for each grid slot, the memory stores the hidden vector of the `RNN`. This memory is used to extend the `RNN` cell, sort of like an `LSTM`.

The loss for training is $\mathcal{L}_{\tau_i, \tau_j} = \sum_k w_k (f(\tau_i, \tau_j) - \exp(-||e_i - e_j||))$, a weighted difference between the similarity metric f and the distance in the embedding space $(e_i - e_j)$. The weight w_k is obtained using the normalized distance matrix, computing pairs of similar and dissimilar trajectories and more fancy stuff.

Notes: map is like google map.

2020 - Trajectory similarity learning with auxiliary supervision and optimal matching

Name: Traj2SimVec

[8] follows the same idea as in [7] which selects some trajectories for pre-training [something], the training samples are divided in three sub-trajectories [because it seems to help learning].

A distance matrix is computed which is used as supervision signal, similar to [7].

2020 - MARC: a robust method for multiple-aspect trajectory classification via space, time, and semantic embeddings

[9] Embeds semantics on the trajectories. Each semantic information (weather, time, type of place) has an encoding, and a weight matrix which transform them into a fixed size vector. The semantic trajectory is fed to an `LSTM`, which encodes the trajectories, having the hidden states used for classification.

2021 - Embedding-Based Similarity Computation for Massive Vehicle Trajectory Data

[10] seems to propose the exact same thing as [7], but with interpolation for de-noising.

2021 - STENet: A hybrid spatio-temporal embedding network for human trajectory forecasting

[11] Focuses on predicting pedestrian trajectories. Uses a `LSTM` with `Convolutional Neural Networks (CNNs)` to embed position features in multiple temporal time-scales. The encoder-decoder structure stack a `CNN` and a graph attention model. The decodes stacks many `LSTMs`.

They give related works on social trajectory learning.

Notes: They point to `Variational Auto-Encoders (VAEs)` for modelling multi-modality and for the generative capabilities.

2021 - A Graph-Based Approach for Trajectory Similarity Computation in Spatial Networks

[12] Propose a **Graph Neural Network (GNN)**-based trajectory embedding. The framework measures trajectory similarities, learns **Points of Interest (PoIs)**, and learns a trajectory embedding.

A trajectory is encoded as the points in a graph map. Then they define a trajectory similarity metric on the **PoI** graph, based on the graph distance between the points and trajectories. An embedding capturing the neighbours and graph trajectory is learned. The **PoI** embeddings and their neighbours are used to learn another embedding using its neighbours information. Finally, **LSTMs** are used to learn the trajectory over the graph embeddings. The loss function minimizes the above defined distance between trajectories and the distance between the two closest trajectories.

2021 - T3S: Effective Representation Learning for Trajectory Similarity Computation

Name: T3S

[13] combines **LSTMs** and attention **NNs** over the grid graph for learning the embedding. Close to [7, 8, 10].

2021 - How meaningful are similarities in deep trajectory representations?

Code on: <https://dbis.ipd.kit.edu/2652.php>

[14] presents a survey and evaluation of **t2vec** [5] and other methods. Seems like **t2vec** with some variations outperform the rest. **t2vec** seems to be stacked **LSTMs**.

Evaluate how changing **t2vec** parameters affect similarity values. **t2vec** seems robust to parameters.

Evaluate **t2vec** against non learning metrics. Seems like associating them lead to better results.

[They DO ignore the whole literature on learning methods?]

Concludes that using **Longest Common SubSequence (LCSS)** and **t2vec** leads to a better trajectory similarity, covering overlap, shape, direction and distance.

Notes: Maybe that should be 4 characteristics to consider.

2022 - Spatio-Temporal Trajectory Similarity Learning in Road Networks

Name: ST2vec

[15] learns a spatio-temporal representation. Two steps, which is based on learning a spatial model, a temporal model and a co-attention fusion module. It is based on a road network, trajectories are sequences of vertex on the road network.

Define the distance of a spatio-temporal trajectory as a weighted sum for a spatio-distance (d_s) and a temporal distance (d_t):

$$d(\tau_i, \tau_j) = \alpha d_s(\tau_i, \tau_j) + (1 - \alpha) d_t(\tau_i, \tau_j) | \alpha \in [0, 1]$$

Later uses **LSTMs** to learn using two strategies, using one **LSTM** for space and another for time, or using one for both.

2022 - Deep Fuzzy Contrast-Set Deviation Point Representation and Trajectory Detection

[16] Grid-map based, contrastive learning.

notes: hard to understand what they are doing here.

2022 - Contrastive Pre-training of Spatial-Temporal Trajectory Embeddings

[17] employs contrastive learning for learning an embedding which retains high-level travel semantics.

Recovering the original trajectory is not a good approach when learning representations with **RNNs** since it fails to capture the high-level information of trajectories. Contrastive learning with noisy augmentation can handle the high-level information while being robust to noise. However data augmentation needs to be well designed.

The positive samples are created with subsampling the query trajectory, while the negative samples come from different trajectories.

Notes: Not sure this is correct, I think the “different trajectories” should be far enough from the query trajectory to be a negative sample.

The encoder stacks a spatio-temporal encoding layer and attention layers. For the first, a learnable encoding of locations is learned (each location leads to a vector) and location and time are passed to a trigonometric vector transformation to compute features which can capture periodic information; those vectors are then summed up. The attention layer is actually 2 stacked attention layers.

2022 - TMN: Trajectory Matching Networks for Predicting Similarity

[18] uses attention to compute intra-trajectory similarities, and then uses a [LSTM](#).

Notes: Comparison ignores many methods.

2022 - TSNE: Trajectory Similarity Network Embedding

[19] uses a pre-defined trajectory measure function to construct a k-NNG (K nearest neighbours graph) and computes the embedding based on the graph.

Notes: Not sure how they compute the embedding from the graph. Seems like the graph representation allows to handle partial similarity and unordered similarity.

2022 - Towards robust trajectory similarity computation: Representation-based spatio-temporal similarity quantification

Name: RSTS

[20] splits the spatio-temporal trajectories into cells, and uses a triplet loss for the learning. It enforces that if the time and space similarities are higher, then the distance in the encoded space must be smaller, and that, in the encoded space, the distance between two trajectories variations (noise and downsampling) must obey the distance of the trajectories.

An embedding is used for the tokens, which are then passed to a [RNN](#) encoder-decoder. The tokens for the embedded are an ID computed by splitting the space-time into cells. The input is grid-cells (gps + time).

Notes: Analysis is poor. Ignores all other works on learning. Seems like there is little innovation besides the loss.

2023 - Spatial-temporal fusion graph framework for trajectory similarity computation

Name: GTS

[21] first learns a point of interest representation on the road network, which is passed to a [GNN](#) for learning neighbours information as embeddings, and then a [LSTM](#) for learning the sequencing.

A symmetric distance between trajectories is defined based on the distance between each point of the trajectories and the other trajectory:

$$d(\tau_1, \tau_2) = \sum_{v \in \tau_1} e^{-d(v, \tau_2)} + \sum_{v \in \tau_2} e^{-d(v, \tau_1)}$$

The time is considered in an extension called ST-LSTM, which adds a time one-hot encoding into the gating functions of the [LSTM](#).

Notes: Comparisons goes as far as traj2SimVec [8].

2023 - GRLSTM: Trajectory Similarity Computation with Graph-Based Residual LSTM

Name: GRLSTM

[22] combines [Knowledge Graph Embedding \(KGE\)](#), [GNN](#) and a multi-layer residual-[LSTM](#). [KGE](#) is used to learn a point and relation embeddings for constructing a graph, which is passed to the [GNN](#) for learning the topology in the point-structure graph. Then the [LSTM](#) is used to learn the embeddings trajectories. Uses two losses: a graph-based loss and a trajectory-base loss.

The input is trajectories in a graph road network. The interesting thing here is that adjacent points in the trajectory may not be adjacent in the graph (due to data loss or lower sample rate).

The stacked [LSTM](#) is augmented with a residual layer for handling the gradient forgetting of traditional [LSTM](#). It is stated that it does not add parameters so it does not affect training time considerably.

Notes: does not really say how the residual function is computed. Similarly to [\[12\]](#) they implement point and trajectory distances.

2023 - Contrastive Trajectory Similarity Learning with Dual-Feature Attention

Name: TrajCL

[\[23\]](#) introduces four trajectory augmentation and a dual feature self-attention encoder, for learning structural and spatial patterns of trajectories. It does not involve any recurrent structure. Instead, it uses a dual self-attention-based trajectory encoder.

Augmentations:

point shifting: adds an offset to the points

point masking randomly removes points from the trajectory

tuncation cuts a prefix, suffix, or both from the trajectory

simplification uses the Douglas–Peucker algorithm which removes non critical points from the trajectories (like points in a straight line).

The augmented trajectories are used to create two trajectory views to learn structural and spatial features. The augmented trajectories are used to compute two trajectory views. The structural features, the map is converted into a grid, and used to create a graph in which the grid locations are the vertices and the trajectory transitions the edges. Then a graph embedding (`node2vec`) is used to learn an embedding. For the spatial features, the angle and length of trajectory segments is computed. Both views are augmented by adding a [\[sketchy\]](#) sine and cosine value to the points to capture position information.

Finally the two views are passed to a two-head self attention module to learn the embeddings.

2023 - Spatio-Temporal Trajectory Similarity Measures: A Comprehensive Survey and Quantitative Study

[\[24\]](#)

Code on: <https://github.com/ZJU-DAILY/TSM>

Present a survey with several methods, and benchmark for evaluating them. Apparently Traj2SimVec [\[8\]](#) is the learning method, which is not grid-based that handles our problem.

2 Trajectory Learning on Robotics

2020 - Controlling Assistive Robots with Learned Latent Actions

[\[25\]](#) Use encoders to learn latent task representations for assistive robot remote controlling. In this setup, [VAEs](#) are used, encoding states into a task representation, the user gives input from a joystick which are decoded together with the latent space representation.

Notes: Seems like [VAE](#) is used straight up with the trajectories. But it is a bit blurry how the actions are being defined or learned (seem pre defined).

2020 - DiversityGAN: Diversity-Aware Vehicle Motion Prediction via Latent Semantic Sampling

[\[26\]](#) extends [Generative Adversarial Network \(GAN\)](#) using a low-dimensional approximate semantic (encoding) which is shaped to capture semantics. Sampling from this space allows to cover semantically distinguish outcomes. The work focuses on predicting vehicle trajectories.

An intermediate layer avoids the need of taxonomy [\[?\]](#) by using metric learning, in which a latent representation is trained to match annotations of high-level labels, and forcing the distance to be large if they represent two distinguish semantic labels. The latent space is trained to match human similarity measures.

Past trajectories and map information are embedded, and their embeddings are passed to an [LSTM](#) whose latent space is divided into a high- and low-level parts. The decoder takes both parts to produce trajectory samples. The trajectory network is a series of fully connected layers that embed a trajectory into a vector

[27] [seems this work uses LSTMs for the embeddings]. The map embedding is a fully connected network that maps polynomial coefficients (quadratic) into an embedding. The encoder is a LSTM, whose hidden states are added a Gaussian noise and passed to a non-linear fully-connected network to compute the high and low-level embedding representation. The high-level embedding part is not correlated with the low-level one, and is trained for learning semantic similarities from the human teacher (they use a hand coded oracle though). The decoder is a LSTM. There is also a discriminator trained for identifying if samples are generated by the architecture or if they are real data.

The loss design incorporates minimal and final displacement losses, a term to enforce the non-correlation between the high and low-level embeddings, and another to enforce that semantically related pairs should also be close in the encoding space.

Sampling is performed using Farthest Point Sampling.

Notes: It is interesting that they added semantics to the network.

2021 - Use of Embedding Spaces for Transferring Robot Skills in Diverse Sensor Settings

[28]

2022 - STL2vec: Signal Temporal Logic Embeddings for Control Synthesis With Recurrent Neural Networks

[29]

2022 - Learning latent actions to control assistive robots

[30]

2022 - Promoting Quality and Diversity in Population-based Reinforcement Learning via Hierarchical Trajectory Space Exploration

[31]

2023 - SIRL: Similarity-Based Implicit Representation Learning

[32]

3 Comparison

4 Challenges

References

- [1] Ilya Sutskever, Oriol Vinyals, and Quoc V Le. Sequence to Sequence Learning with Neural Networks. In Z. Ghahramani, M. Welling, C. Cortes, N. Lawrence, and K.Q. Weinberger, editors, *Advances in Neural Information Processing Systems*, volume 27. Curran Associates, Inc., 2014.
- [2] Qiang Gao, Fan Zhou, Kunpeng Zhang, Goce Trajcevski, Xucheng Luo, and Fengli Zhang. Identifying Human Mobility via Trajectory Embeddings. In *Proceedings of the 26th International Joint Conference on Artificial Intelligence, IJCAI'17*, page 1689–1695. AAAI Press, 2017.
- [3] John Co-Reyes, YuXuan Liu, Abhishek Gupta, Benjamin Eysenbach, Pieter Abbeel, and Sergey Levine. Self-Consistent Trajectory Autoencoder: Hierarchical Reinforcement Learning with Trajectory Embeddings. In Jennifer Dy and Andreas Krause, editors, *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pages 1009–1018. PMLR, 10–15 Jul 2018.
- [4] Li Song, Ruijia Wang, Ding Xiao, Xiaotian Han, Yanan Cai, and Chuan Shi. Anomalous Trajectory Detection Using Recurrent Neural Network. In Guojun Gan, Bohan Li, Xue Li, and Shuliang Wang, editors, *Advanced Data Mining and Applications*, pages 263–277, Cham, 2018. Springer International Publishing.
- [5] Xiucheng Li, Kaiqi Zhao, Gao Cong, Christian S. Jensen, and Wei Wei. Deep Representation Learning for Trajectory Similarity Computation. In *2018 IEEE 34th International Conference on Data Engineering (ICDE)*, pages 617–628, 2018.
- [6] Tao-Yang Fu and Wang-Chien Lee. Trembr: Exploring Road Networks for Trajectory Representation Learning. *ACM Trans. Intell. Syst. Technol.*, 11(1), feb 2020.
- [7] Di Yao, Gao Cong, Chao Zhang, and Jingping Bi. Computing Trajectory Similarity in Linear Time: A Generic Seed-Guided Neural Metric Learning Approach. In *2019 IEEE 35th International Conference on Data Engineering (ICDE)*, pages 1358–1369, 2019.
- [8] Hanyuan ZHANG, Xingyu ZHANG, Qize JIANG, Baihua ZHENG, Zhenbang SUN, Weiwei SUN, and Changhu WANG. Trajectory similarity learning with auxiliary supervision and optimal matching. In *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence*, pages 3209–3215. Research Collection School Of Information Systems, 11–17 Jul 2020.
- [9] Lucas May Petry, Camila Leite Da Silva, Andrea Esuli, Chiara Renso, and Vania Bogorny. MARC: a robust method for multiple-aspect trajectory classification via space, time, and semantic embeddings. *International Journal of Geographical Information Science*, 34(7):1428–1450, 2020.
- [10] Yuanyi Chen, Peng Yu, Wenwang Chen, Zengwei Zheng, and Minyi Guo. Embedding-Based Similarity Computation for Massive Vehicle Trajectory Data. *IEEE Internet of Things Journal*, 9(6):4650–4660, 2022.
- [11] Bo Zhang, Chengzhi Yuan, Tao Wang, and Hongbo Liu. STENet: A hybrid spatio-temporal embedding network for human trajectory forecasting. *Engineering Applications of Artificial Intelligence*, 106:104487, 2021.
- [12] Peng Han, Jin Wang, Di Yao, Shuo Shang, and Xiangliang Zhang. A Graph-Based Approach for Trajectory Similarity Computation in Spatial Networks. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining, KDD '21*, page 556–564, New York, NY, USA, 2021. Association for Computing Machinery.
- [13] Peilun Yang, Hanchen Wang, Ying Zhang, Lu Qin, Wenjie Zhang, and Xuemin Lin. T3S: Effective Representation Learning for Trajectory Similarity Computation. In *2021 IEEE 37th International Conference on Data Engineering (ICDE)*, pages 2183–2188, 2021.
- [14] Saeed Taghizadeh, Abel Elekes, Martin Schäler, and Klemens Böhm. How meaningful are similarities in deep trajectory representations? *Information Systems*, 98:101452, 2021.
- [15] Ziquan Fang, Yuntao Du, Xinjun Zhu, Danlei Hu, Lu Chen, Yunjun Gao, and Christian S. Jensen. Spatio-Temporal Trajectory Similarity Learning in Road Networks. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, KDD '22*, page 347–356, New York, NY, USA, 2022. Association for Computing Machinery.

- [16] Usman Ahmed, Jerry Chun-Wei Lin, and Gautam Srivastava. Deep Fuzzy Contrast-Set Deviation Point Representation and Trajectory Detection. *IEEE Transactions on Fuzzy Systems*, 31(2):571–581, 2023.
- [17] Yan Lin, Huaiyu Wan, Shengnan Guo, and Youfang Lin. Contrastive Pre-training of Spatial-Temporal Trajectory Embeddings, 2022.
- [18] Peilun Yang, Hanchen Wang, Defu Lian, Ying Zhang, Lu Qin, and Wenjie Zhang. TMN: Trajectory Matching Networks for Predicting Similarity. In *2022 IEEE 38th International Conference on Data Engineering (ICDE)*, pages 1700–1713, 2022.
- [19] Jiaxin Ding, Bowen Zhang, Xinbing Wang, and Chenghu Zhou. TSNE: Trajectory Similarity Network Embedding. In *Proceedings of the 30th International Conference on Advances in Geographic Information Systems, SIGSPATIAL '22*, New York, NY, USA, 2022. Association for Computing Machinery.
- [20] Ziwen Chen, Ke Li, Silin Zhou, Lisi Chen, and Shuo Shang. Towards robust trajectory similarity computation: Representation-based spatio-temporal similarity quantification. *World Wide Web*, pages 1–24, 2022.
- [21] Silin Zhou, Peng Han, Di Yao, Lisi Chen, and Xiangliang Zhang. Spatial-temporal fusion graph framework for trajectory similarity computation. *World Wide Web*, 26(4):1501–1523, 2023.
- [22] Silin Zhou, Jing Li, Hao Wang, Shuo Shang, and Peng Han. GRLSTM: Trajectory Similarity Computation with Graph-Based Residual LSTM. *Proceedings of the AAAI Conference on Artificial Intelligence*, 37(4):4972–4980, Jun. 2023.
- [23] Yanchuan Chang, Jianzhong Qi, Yuxuan Liang, and Egemen Tanin. Contrastive Trajectory Similarity Learning with Dual-Feature Attention. In *2023 IEEE 39th International Conference on Data Engineering (ICDE)*, pages 2933–2945, 2023.
- [24] Danlei Hu, Lu Chen, Hanxi Fang, Ziquan Fang, Tianyi Li, and Yunjun Gao. Spatio-Temporal Trajectory Similarity Measures: A Comprehensive Survey and Quantitative Study, 2023.
- [25] Dylan P. Losey, Krishnan Srinivasan, Ajay Mandlekar, Animesh Garg, and Dorsa Sadigh. Controlling Assistive Robots with Learned Latent Actions. In *2020 IEEE International Conference on Robotics and Automation (ICRA)*, pages 378–384, 2020.
- [26] Xin Huang, Stephen G. McGill, Jonathan A. DeCastro, Luke Fletcher, John J. Leonard, Brian C. Williams, and Guy Rosman. DiversityGAN: Diversity-Aware Vehicle Motion Prediction via Latent Semantic Sampling. *IEEE Robotics and Automation Letters*, 5(4):5089–5096, 2020.
- [27] Alexandre Alahi, Kratarth Goel, Vignesh Ramanathan, Alexandre Robicquet, Li Fei-Fei, and Silvio Savarese. Social lstm: Human trajectory prediction in crowded spaces. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2016.
- [28] Kazushi Ninomiya, Masakazu Hirokawa, and Kenji Suzuki. Use of Embedding Spaces for Transferring Robot Skills in Diverse Sensor Settings. In *2021 20th IEEE International Conference on Machine Learning and Applications (ICMLA)*, pages 547–551, 2021.
- [29] Wataru Hashimoto, Kazumune Hashimoto, and Shigemasa Takai. STL2vec: Signal Temporal Logic Embeddings for Control Synthesis With Recurrent Neural Networks. *IEEE Robotics and Automation Letters*, 7(2):5246–5253, 2022.
- [30] Dylan P Losey, Hong Jun Jeon, Mengxi Li, Krishnan Srinivasan, Ajay Mandlekar, Animesh Garg, Jeannette Bohg, and Dorsa Sadigh. Learning latent actions to control assistive robots. *Autonomous robots*, 46(1):115–147, 2022.
- [31] Jiayu Miao, Tianze Zhou, Kun Shao, Ming Zhou, Weinan Zhang, Jianye Hao, Yong Yu, and Jun Wang. Promoting Quality and Diversity in Population-based Reinforcement Learning via Hierarchical Trajectory Space Exploration. In *2022 International Conference on Robotics and Automation (ICRA)*, pages 7544–7550, 2022.
- [32] Andreea Bobu, Yi Liu, Rohin Shah, Daniel S. Brown, and Anca D. Dragan. SIRL: Similarity-Based Implicit Representation Learning. In *Proceedings of the 2023 ACM/IEEE International Conference on Human-Robot Interaction, HRI '23*, page 565–574, New York, NY, USA, 2023. Association for Computing Machinery.